



In the name of Allah

# Physics-Informed Neural Network: Notes

By:

**Ehsan Ghaderi**

Adviser:


**Dr. Bijarchi**

October 18,2023

# Latin Hypercube Sampling (LHS)

The idea behind one-dimensional latin hypercube sampling is simple:

Divide a given **CDF** into  $n$  different regions and randomly choose one value from each region to obtain a sample of size  $n$ .




The benefit of this approach is that it ensures that at least one value from each region is included in the sample.

# Latin Hypercube Sampling (LHS)

We can easily extend the idea of one-dimensional latin hypercube sampling into two dimensions as well.

For two variables,  $x$  and  $y$ , we can divide the sample space of each variable into  $n$  evenly spaced regions and pick a random sample from each sample space to obtain random values across two dimensions.



# LHS

```
def lhs(n, samples=None, criterion=None, iterations=None):
```

```
    """
```

```
    Generate a latin-hypercube design
```

```
    Parameters
```

```
    -----
```

```
    n : int
```

```
        The number of factors to generate samples for
```

```
    Optional
```

```
    -----
```

```
    samples : int
```

```
        The number of samples to generate for each factor (Default: n)
```

```
    criterion : str
```

```
        Allowable values are "center" or "c", "maximin" or "m",
```

```
        "centermaximin" or "cm", and "correlation" or "corr". If no value  
        given, the design is simply randomized.
```

```
    iterations : int
```

```
        The number of iterations in the maximin and correlations algorithms  
        (Default: 5).
```

```
>>> lhs(2, samples=5, criterion='center')  
array([[ 0.3,  0.5],  
       [ 0.7,  0.9],  
       [ 0.1,  0.3],  
       [ 0.9,  0.1],  
       [ 0.5,  0.7]])
```

# Initialization

- Xavier Normal initialization

$$W \sim N(0, Var(W))$$

*Handwritten notes:* A red circle around  $N$  with an arrow pointing to "layer-1". A blue circle around  $Var(W)$  with an arrow pointing to "std".

$$Var(W) = \sqrt{\frac{2}{n_{in} + n_{out}}}$$

*Handwritten notes:* "Standard deviation : 2" written in blue above the square root.

- Xavier Uniform initialization

$$W \sim U\left(-\sqrt{\frac{6}{n_{in} + n_{out}}}, +\sqrt{\frac{6}{n_{in} + n_{out}}}\right)$$

*Handwritten notes:* A blue circle around  $n_{out}$  in the denominator.

- He Normal initialization

$$W \sim N(0, Var(W))$$

$$Var(W) = \sqrt{\frac{2}{n_{in}}}$$

- He Uniform initialization

$$W \sim U\left(-\sqrt{\frac{6}{n_{in}}}, +\sqrt{\frac{6}{n_{in}}}\right)$$

# Xavier initialization

```
def xavier_normal_(tensor: Tensor, gain: float = 1.) -> Tensor:
    r"""Fills the input `Tensor` with values according to the method
    described in `Understanding the difficulty of training deep feedforward
    neural networks` - Glorot, X. & Bengio, Y. (2010), using a normal
    distribution. The resulting tensor will have values sampled from
    :math:\mathcal{N}(0, \text{std}^2)` where

    .. math::
        \text{std} = \text{gain} \times \sqrt{\frac{2}{\text{fan\_in} + \text{fan\_out}}}
```

Also known as Glorot initialization.

Args:

tensor: an n-dimensional `torch.Tensor`  
gain: an optional scaling factor

Examples:

```
>>> w = torch.empty(3, 5)
>>> nn.init.xavier_normal_(w)
```

Weights are taken from the normal distribution whose mean is equal to zero and standard deviation equal to  $\sqrt{2/\text{fan\_in} + \text{fan\_out}}$ ; fan\_out is the number of outputs from the neuron.

```
"""
```

```
fan_in, fan_out = _calculate_fan_in_and_fan_out(tensor)
std = gain * math.sqrt(2.0 / float(fan_in + fan_out))
```

```
return _no_grad_normal_(tensor, 0., std)
```



## Optimization algorithms

L-BFGS	ADAM
$x_{k+1} = x_k + \alpha_k d_k$ $s_k = x_{k+1} - x_k$ $y_k = g_{k+1} - g_k$ $\alpha_k = \frac{s_k^T y_k}{y_k^T H_k y_k}$ $x_{k+1} = x_k + \alpha_k d_k$	$m_k = \beta_1 m_{k-1} + (1 - \beta_1) g_k$ $v_k = \beta_2 v_{k-1} + (1 - \beta_2) g_k^2$ $m'_k = \frac{m_k}{1 - \beta_1^k}$ $v'_k = \frac{v_k}{1 - \beta_2^k}$ $x_{k+1} = x_k - \frac{\eta}{\sqrt{v'_k} + \epsilon} m'_k$

# Adam

$$m_w^{t+1} = \beta_1 m_w^t + (1 - \beta_1) \nabla_w L^t$$

$$v_w^{t+1} = \beta_2 v_w^t + (1 - \beta_2) (\nabla_w L^t)^2$$

$$\hat{m}_w = \frac{m_w^{t+1}}{1 - \beta_1^t}$$

$$\hat{v}_w = \frac{v_w^{t+1}}{1 - \beta_2^t}$$

$$w^{t+1} = w^t - \eta \frac{\hat{m}_w}{\sqrt{\hat{v}_w} + \epsilon}$$

where  $w(t)$  are the model parameters,  $L^t$  is the loss function,  $t$  is the current training iteration,  $\beta_1$  and  $\beta_2$  are the forgetting factors for gradients and second-order gradient moments, respectively.



# Adam

```
def __init__(self, params, lr=1e-3, betas=(0.9, 0.999), eps=1e-8,  
             weight_decay=0, amsgrad=False):
```

Args:

params (iterable): iterable of parameters to optimize or dicts defining parameter groups

lr (float, optional): learning rate (default: 1e-3)

betas (Tuple[float, float], optional): coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))

eps (float, optional): term added to the denominator to improve numerical stability (default: 1e-8)

weight\_decay (float, optional): weight decay (L2 penalty) (default: 0)

amsgrad (boolean, optional): whether to use the AMSGrad variant of this algorithm from the paper `'On the Convergence of Adam and Beyond'` (default: False)

---

**Algorithm 1** L-BFGS

---

**for**  $i = 1, 2, \dots, n$  **do**

Obtain a direction  $P_k$  by solving  $B_k P_k = -\nabla f(X_k)$

Perform a one-dimensional optimization to find an acceptable stepsize  $\alpha_k$  in the direction found in the first step. In an exact line search,  $\alpha_k = \operatorname{argmin} f(X_k + \alpha P_k)$ .

In practice, an inexact line search usually suffices, with acceptable  $\alpha_k$  satisfying the Wolfe conditions.

Set  $S_k = \alpha_k P_k$  and update  $X_{k+1} = X_k + S_k$

$$y_k = \nabla f(X_{k+1}) - \nabla f(X_k)$$

$$B_{k+1} = B_k + \frac{Y_k Y_k^T}{s_k Y_k^T} + \frac{B_k s_k s_k^T B_k^T}{B_k s_k s_k^T}$$

**end for**

---

where  $X_k$  are the model parameters,  $f$  is the loss function,  $B_k$  is an approximation of the Hessian matrix, and  $S_k$  and  $Y_k$  are the differences between the current and previous values of  $X$  and  $\nabla f(X)$ , respectively. The L-BFGS and ADAM optimization algorithms are both commonly

# L-BFGS

```
def __init__(self,  
    params,  
    lr=1,  
    max_iter=20,  
    max_eval=None,  
    tolerance_grad=1e-7,  
    tolerance_change=1e-9,  
    history_size=100,  
    line_search_fn=None):
```

Args:

lr (float): learning rate (default: 1)  
max\_iter (int): maximal number of iterations per optimization step  
(default: 20)  
max\_eval (int): maximal number of function evaluations per optimization  
step (default: max\_iter \* 1.25).  
tolerance\_grad (float): termination tolerance on first order optimality  
(default: 1e-5).  
tolerance\_change (float): termination tolerance on function  
value/parameter changes (default: 1e-9).  
history\_size (int): update history size (default: 100).  
line\_search\_fn (str): either 'strong\_wolfe' or None (default: None).

"""